

Navigation

October 4, 2019

1 Navigation

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [1]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 2.0.
```

The environment is already saved in the Workspace and can be accessed at the file path provided below. Please run the next code cell without making any changes.

```
In [2]: from unityagents import UnityEnvironment
import numpy as np

# please do not modify the line below
env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")
```

```
INFO:unityagents:
```

```
'Academy' started successfully!
```

```
Unity Academy name: Academy
```

```
Number of Brains: 1
```

```
Number of External Brains : 1
```

```
Lesson number : 0
```

```
Reset Parameters :
```

```
Unity brain name: BananaBrain
```

```
Number of Visual Observations (per agent): 0
```

```
Vector Observation space type: continuous
```

```
Vector Observation space size (per agent): 37
```

```

Number of stacked Vector Observation: 1
Vector Action space type: discrete
Vector Action space size (per agent): 4
Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [3]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

```

1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [4]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents in the environment
        print('Number of agents:', len(env_info.agents))

        # number of actions
        action_size = brain.vector_action_space_size
        print('Number of actions:', action_size)

        # examine the state space
        state = env_info.vector_observations[0]
        print('States look like:', state)
        state_size = len(state)
        print('States have length:', state_size)

```

```

Number of agents: 1
Number of actions: 4
States look like: [ 1.          0.          0.          0.          0.84408134  0.          0.
 1.          0.          0.0748472  0.          1.          0.          0.
 0.25755      1.          0.          0.          0.          0.74177343
 0.          1.          0.          0.          0.25854847  0.          0.
 1.          0.          0.09355672  0.          1.          0.          0.
 0.31969345  0.          0.          ]
States have length: 37

```

1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that in this coding environment, you will not be able to watch the agent while it is training, and you should set `train_mode=True` to restart the environment.

```
In [5]: env_info = env.reset(train_mode=True)[brain_name] # reset the environment
        state = env_info.vector_observations[0]           # get the current state
        score = 0                                         # initialize the score
        while True:
            action = np.random.randint(action_size)      # select an action
            env_info = env.step(action)[brain_name]      # send the action to the environment
            next_state = env_info.vector_observations[0]  # get the next state
            reward = env_info.rewards[0]                 # get the reward
            done = env_info.local_done[0]                # see if episode has finished
            score += reward                               # update the score
            state = next_state                           # roll over the state to next time step
            if done:                                     # exit loop if episode finished
                break

        print("Score: {}".format(score))

Score: 0.0
```

When finished, you can close the environment.

```
In [6]: # env.close()
```

1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agent while it is training. However, *after training the agent*, you can download the saved model weights to watch the agent on your own machine!

1.0.5 My Implementation

```
In [7]: import random
        import torch
        from collections import deque
        import time
        import matplotlib.pyplot as plt
        %matplotlib inline
```

First, import your DQN Agent

```
In [8]: from dqn_agent import Agent
```

```
# lets set the states and action for the agent (Information is printed above)
agent = Agent(state_size=37, action_size=4, seed=50)
```

Deep Q-Learning. Parameter definitions

- n_episodes (int): maximum number of training episodes
- max_t (int): maximum number of timesteps per episode
- eps_start (float): starting value of epsilon, for epsilon-greedy action selection
- eps_end (float): minimum value of epsilon
- eps_decay (float): multiplicative factor (per episode) for decreasing epsilon

Parameters	Values
n_episodes	2000
max_t	1000
eps_start	1
eps_end	0.01
eps_decay	0.9995

Parameter Values Chosen

Learning Algorithm - Deep Q-Learning Algorithm Start with an initial action-value $Q_0(s, a)$ for all states and action s, a . $Q(\text{terminal state}) = 0$

for $episode = 1, 2, \dots$ until number of episodes: - Get initial state s from the environment - for $n = 1, 2, \dots$ until convergence: - - sample probabilities of actions a - - execute action in the environment at state s , get reward and next state s' - - if s' is terminal: - - - target = $R(s, a, s')$ - - - sample new initial state s' - - else: - - - target = $R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$ - - $Q_{k+1}(s, a) = (1 - \alpha) Q_k(s, a) + [\text{target}]$ - - ss'

Model Architectures - A neural network which consists of two fully connected layers with 64 units each together with an output layer

```
In [9]: def dqn(n_episodes=2000, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.9995):

    scores = [] # Initialize list to save scores from each episode
    scores_window = deque(maxlen=100) # Last 100 scores
    eps = eps_start # initialize epsilon
    for i_episode in range(1, n_episodes+1):
        env_space = env.reset(train_mode=True)[brain_name] # reset the state environment
        state = env_space.vector_observations[0] # get the current state

        score = 0
```

```

for t in range( max_t):
    action = agent.act(state, eps)    # get a probability of states
    env_space = env.step(action)[brain_name]    # take action in the environment
    reward = env_space.rewards[0]        # get reward
    next_state = env_space.vector_observations[0]    # get next state
    done = env_space.local_done[0]    # check to see if current episode has finished
    agent.step(state, action, reward, next_state, done)
    score += reward    #update score
    state = next_state
    if done:
        break

scores_window.append(score)    # save most recent score
scores.append(score)
eps = max(eps_end, eps_decay*eps) # decrease epsilon
print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
if i_episode % 100 == 0:
    print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
if np.mean(scores_window) >= 13.0:
    print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
    torch.save(agent.qnetwork_local.state_dict(), 'checkpoint.pth')
    break
return scores

```

```

In [10]: # Train the agent
start_time = time.time() # Monitor Training Time
scores = dqn()
print("\nTotal Training time = {:.1f} min".format((time.time()-start_time)/60))

```

```

Episode 100      Average Score: 0.89
Episode 200      Average Score: 4.53
Episode 300      Average Score: 7.25
Episode 400      Average Score: 10.75
Episode 469      Average Score: 13.00
Environment solved in 369 episodes!      Average Score: 13.00


```

```
Total Training time = 9.8 min
```

```

In [11]: # plot the scores
fig = plt.figure(figsize=(8,8))
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.title('Score (Rewards)')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.grid(True)
plt.show()

```



1.0.6 Near Future Improvement and Exploration to Tryout

- Train the agent using raw screen pixels as input [Link to Implementation](#)
- Improve performance using Prioritized Experience Replay. [Prioritized Experience Replay using a special data structure Sum Tree](#)
- Other Improvements in Deep Q Learning [Link to Implementation](#)

In []: